

what does(n't) work for chatbots

Oleksandr Khrypilyvenko
m3oucat@gmail.com

#####

#####

have you tried fixing your device

#####

with a screw driver? | 10W 50K vocab


#####

#####


try screw driver | 3W 4K vocab

####

~~Garbage in → Garbage out~~



NNs suffer from imbalanced datasets



I'd like to refund

i don't know

my unit is broken

i don't know

do you work?


i don't know


can i see your manager?

i don't know

whatever question?

i don't know





data cleaning preprocessing

vocabulary reduction

data separation

labeling



clustering

dataset balancing




vocabulary reduction

frequency, lemmatization, spelling




dataset balancing

reduce bigger, increase smaller



labeling

intents, slots, rewards, start/end of answer



data separation

conversational, query, task completion




cluster separated data

for typical problems



NLU

o=question



**State Machine
(SM)**

$s, \pi : S \rightarrow A$



NLG

a=answer

NLU

selects a model to use from

{conversational, query, task completion}

intent detection & slot values <https://dialogflow.com/>

"wanna pizza margarita @ work, pay cash"

order_pizza(loc=_, pay=cash)

SM

makes you able to define policies

"wanna pizza margarita @ work, pay cash"

o(bservervation) = `order_pizza(loc=_, pay=cash)`

H(istory): work: loc = 49,72

S(tate) = `order_pizza(loc=49,72, pay=cash)`

π (olicy): $S \rightarrow A$ (ction space)

π (`order_pizza(loc=*, pay=*)`) =

= `a_deliver(order(loc=loc, pay=pay))` = `a(ction)`

NLG

runs on state machine

depends on policy

templates(user-friendly answers)

```
a_greet = random.choice(["Hi", "Hello", "Greetings"])
```

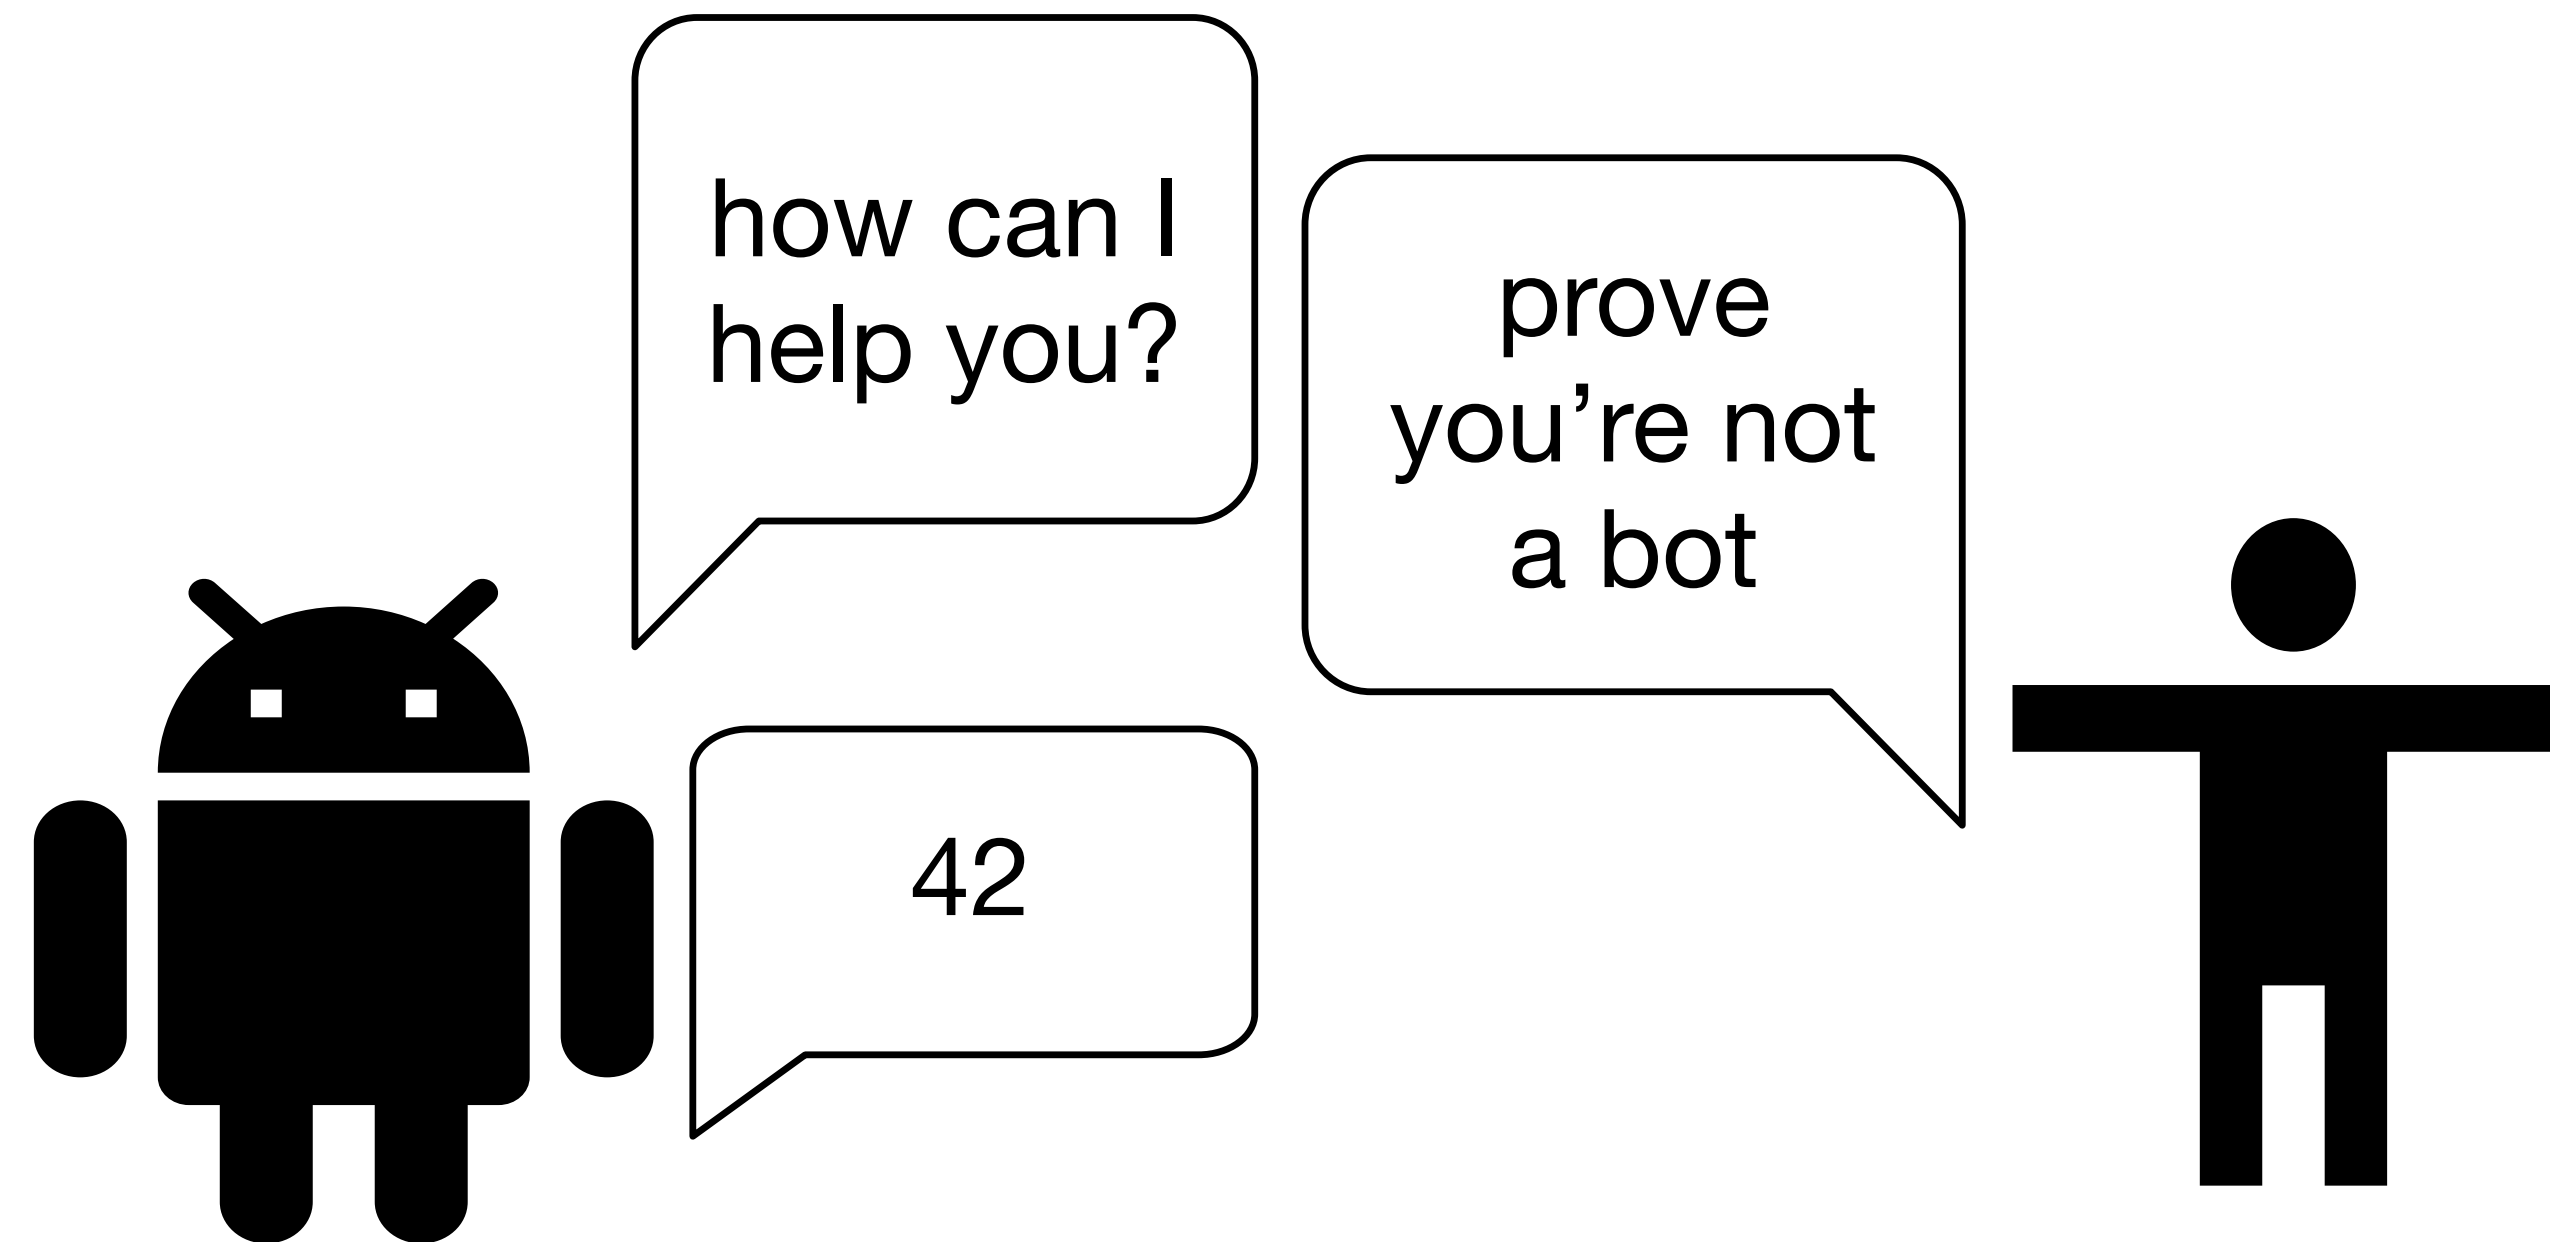


conversational

task completion(TC)

text query(TQ)

conversational sustain a chat



covers a lot of existing text w/o labeling
NLU, SM, NLG, TQ: all in one
suffers from imbalanced datasets
hard to train
hard to manage NERs

conversational

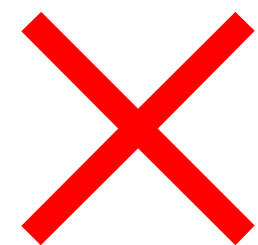
plain seq2seq

seq2seq + attention

transformer

BIG models/ensembles

ranging

 plain seq2seq
(uni/bi directional LSTMs)

lots of samples

simple to grasp

very high probability to stuck in

"i don't know" | dirty data

✓ seq2seq + attention

official version is easy to modify

if cooked properly, performs the best

beam search out-of-the box

MMI may be easily added

works(better) when trained on one to many

takes more time to train than transformer

<https://github.com/tensorflow/nmt>

<https://arxiv.org/pdf/1510.03055.pdf> MMI

<https://www.youtube.com/watch?v=RLWuzLLSlgw> beam search

i can't use this **device** in my life

you can use the browse by **device** type make
search engine on our home page

i searched there and it says it should work

if you have any question later let us know
/common answer | unrelated context/

i would like to **refund**

i will let them know to **refund** you the **refund**

may i talk to your **manager**

i will let **her** know to send it to you

you're pissing me off

yes

why you're doing this?

i do not know if it is possible to be

to be what?

you can also use the browse

i made browsing and it assures me i should use my charger

you can check our youtube channel to see if there is anything
you could do

/from time to time it gives helpful answers because trained
on good data/

i just watched your channel and it says i'm charging it in wrong way

i am sorry i am glad i can help

what are you sorry of? you haven't helped me!

i am not sure what you mean by more question i will have to check on forum for that device

thank you a lot for keeping trying

i will let her know to send you the document

you just ensured me that you'll look for a solution on forum

i will have to check on forum coupon code

no need a coupon, need my device working!

you can use the search engine on our home page to link you to the manual for your device

nice chat sample

my unit doesn't work

what firmware version are you use

i'm using firmware version 2.1

you can try update firmware if you would

ok, I'll update it, thanks

yes correct

✗ transformer

available on github

feedforward

quick to train

does worse than nmt

official version is not so easy to customize

too generic answers(only attention)

my ubuntu doesn't work

i would say that is a bit of a of
a bit of what?

i would say bit

can you tell me something meaningful?

i do not know

i can't get online on my laptop

what is the problem

i can't install software

what is the error

the error is: exception case

i would try the alternate eol of the ubuntu

BIG models, ensembles

if you've got lucky, you get diverse answers

require too much resources

not so easy to modify, even run

high chance "I don't know" | YOUR data

require labeled data

ranging (concept)

$$a = \operatorname{argmin} D(a_{user}, \forall a \in A)$$

$$q = \operatorname{argmin} D(q_{user}, \forall q \in Q)$$

$$\pi = f: Q \rightarrow A$$

$$a = \max_{score} a \mid q$$

simple to implement

easy to debug

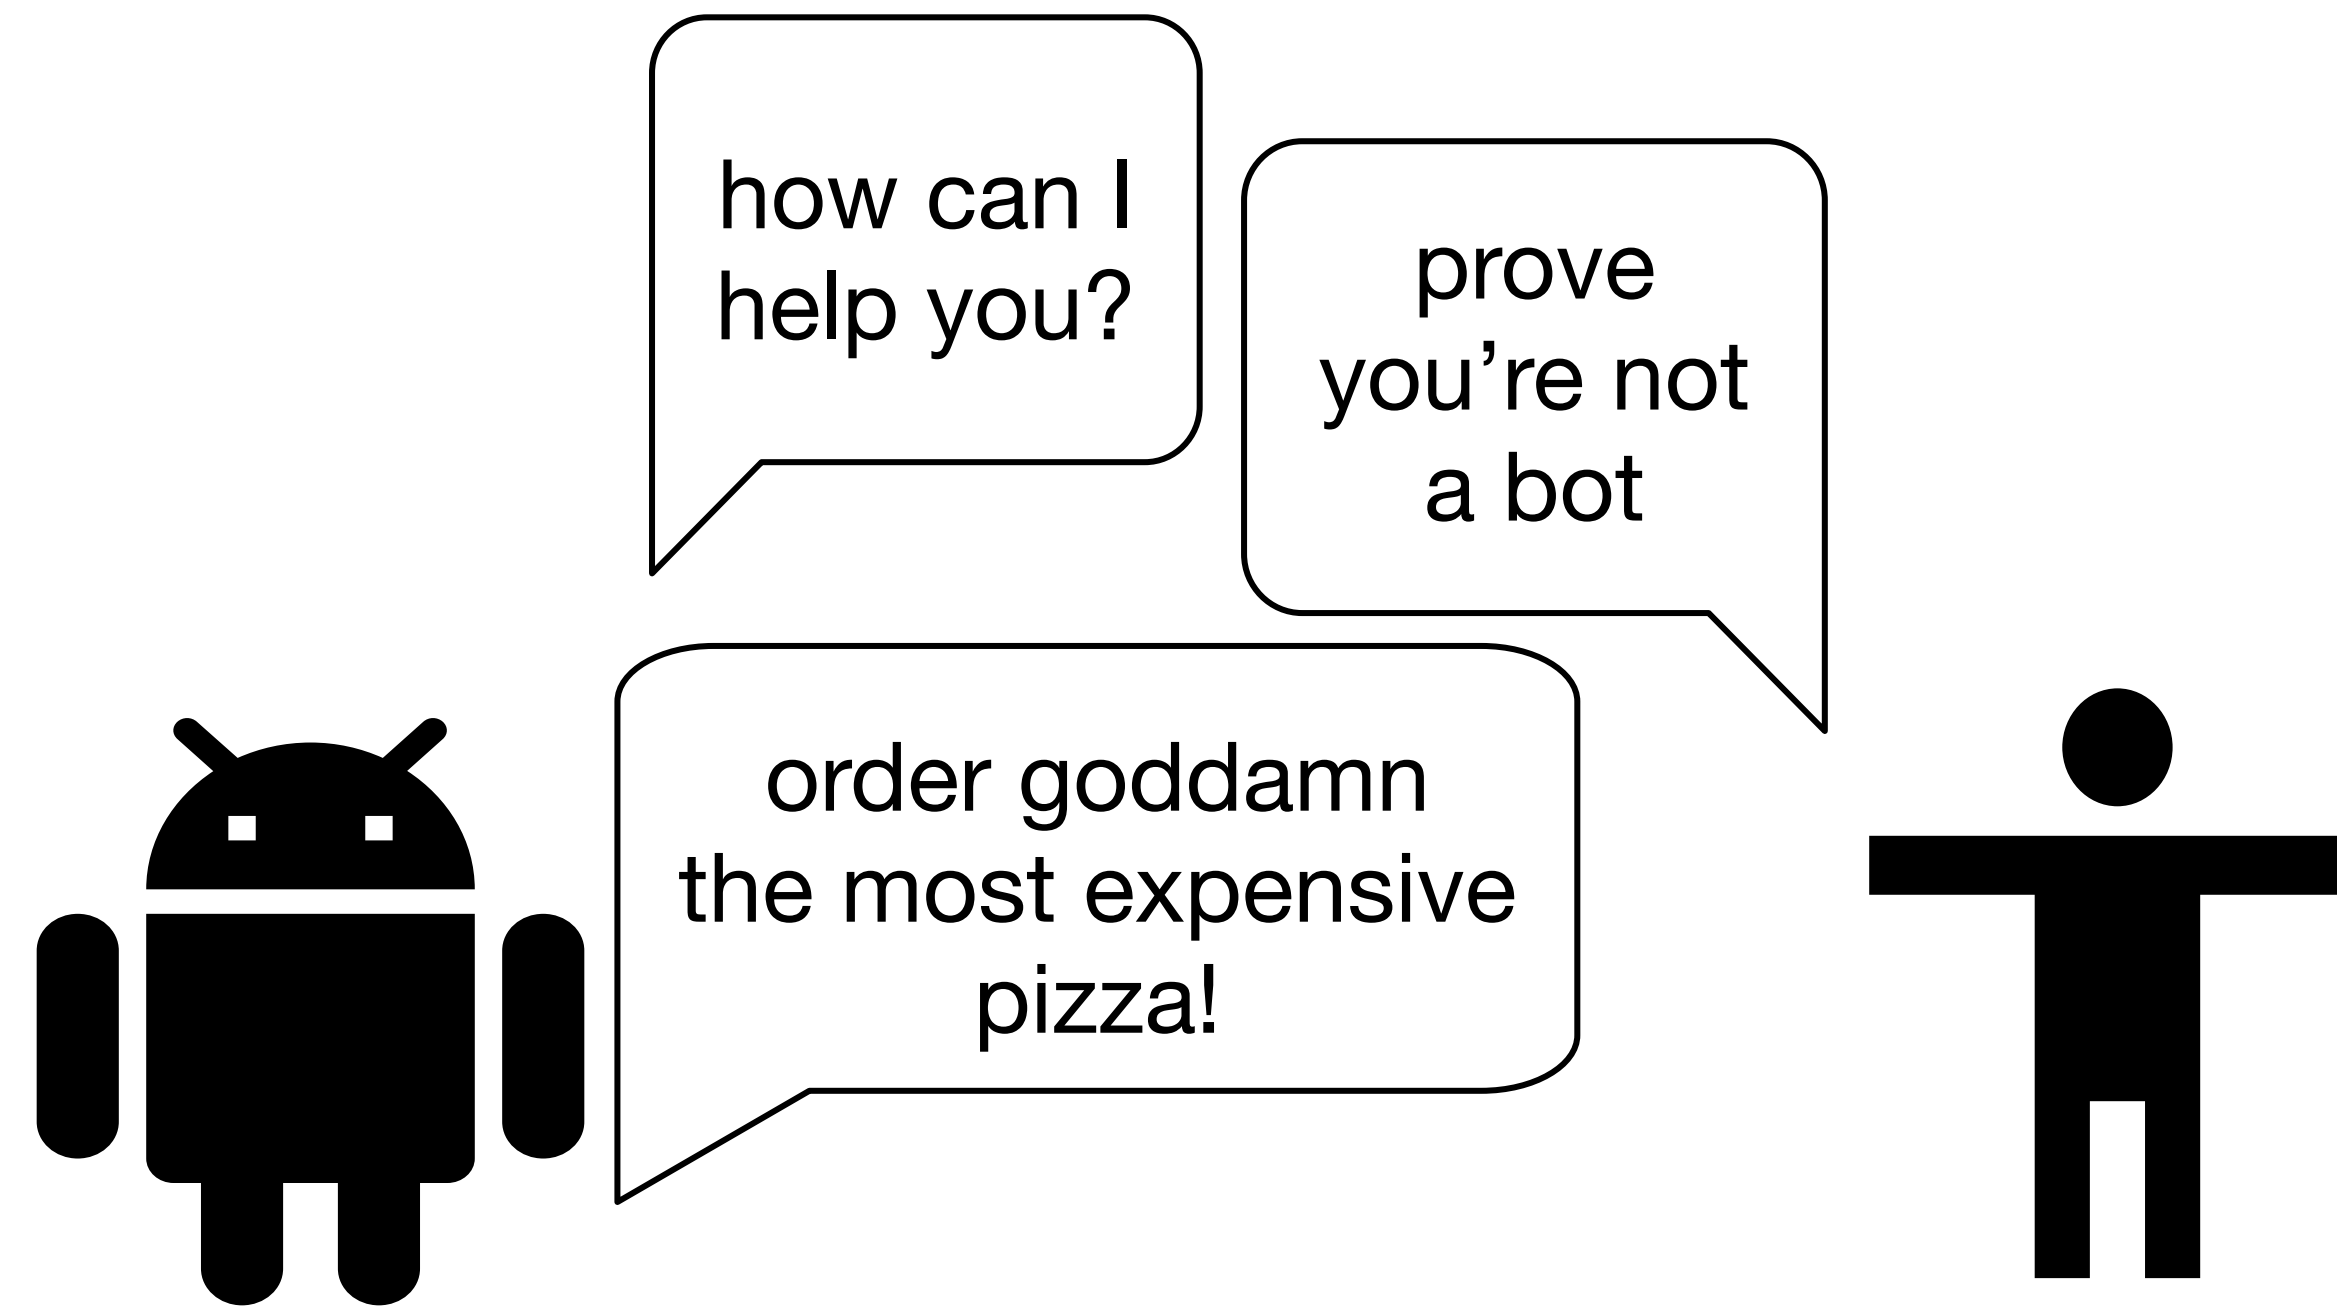
works for TQ(text query)

heavily depends on encoding quality

(huge labeled datasets)

TC

have to complete a task
or conclude there's no way you can do it



tractable, easy to debug & modify
requires lots of manual work
most used approaches require NLU

✗ (RL) in discrete space

$$\begin{array}{c} LSH \\ f : \underbrace{V^\ell}_{\text{pretrained NN}} \rightarrow \mathbb{R}^n \xrightarrow{\quad} \underbrace{\{0,1\}^m}_{\pi : S \rightarrow A} \rightarrow V^\ell \\ \text{word/sentence} \rightarrow \text{embedding} \end{array}$$

hard to keep LSH as non-linear
homomorphism(metric changes)

pretrained NN = skipthoughts, BiMPM...


Reward depends on how successfully you complete a task

? (RL) in continuous space

state definition

- We may say, we simply learn policy $\pi : S \rightarrow A$
- q/a = all successive user's/agent's sentences(utterance) before agent's/(user's) ones respectively
- $a = a$
- $s = q(aq)(aq)(aq)\dots$ | all $(q\dots q)$ so far

$$a = \arg \max_Q \frac{\partial Q(s, a)}{\partial a}$$



$$(q \dots), (a) \in vocab \quad s \in \mathbb{R}^m, a \in \mathbb{R}^n, Q \in \mathbb{R}$$

the problem even worse when a is a sentence-wise vector

$$a \in \mathbb{R}^n!$$

getting back to discrete V space is non-trivial

we have to train another network

$$a \in \mathbb{R}^n \rightarrow a \in V^\ell, \ell = \text{sen len in words}$$

✓ Forms

interpretable

allows RL in discrete space(converges)

datasets are optional (coldstart)

lots of manual work

need NLU

need to define default policy

if state defined wrong, pisses off users

SM test sample

```
TEST_CASE = [  
    [  
        "Hey, buddy, I have nothing to do. I'd like to talk to someone",  
        [  
            ("NO_INTENTS_DETECTED", dict()),  
        ],  
    ],  
    [  
        "Could you give me a pepperoni and one more",  
        [  
            ("specify_pizza", {"pizza_type": "pepperoni", "quantity": 1}),  
            ("order_pizza", dict()),  
        ],  
    ],  
    ...,  
    [  
        "yep, that's all",  
        [  
            ("confirm", dict())  
        ],  
    ],  
]
```

Hey, buddy, I have nothing to do. I'd like to talk to someone

Hello

What kind of pizza do you want?

you can select pre-set pizzaz

or you can make you own pizza from scratch from the following ingredients:

cheddar salami mozzarella tomatoes sweet red pepper pineapple chicken bacon
chili red beef

preset pizza types are: margarita pepperoni chicken vegetarian calzone

Could you give me a pepperoni and one more

What kind of pizza do you want?

you can select pre-set pizzaz

or you can make you own pizza from scratch from the following ingredients:

cheddar salami mozzarella tomatoes sweet red pepper pineapple chicken bacon
chili red beef

preset pizza types are: margarita pepperoni chicken vegetarian calzone

which address deliver to?

delivery time?

which payment type do you prefer?

...

Would you like to add something else to your order?

yep, that's all

```
{'delivery_address': '80 banana st', 'delivery_time': '14:00', 'payment_type':  
'cash', 'cart': [('pizza', {'pizza_type': 'pepperoni', 'quantity': 1}),  
('custom_pizza', {'quantity': 1, 'ingredients': {'cheese': 1}})]}
```


S definition

question:

`words = sentence.lower().split()`

'cat' and 'meow' in words = s_{cat}^1

'cat' and 'scratch' in words = s_{cat}^2

'dog' and 'barks' in words = s_{dog}^1

(this something like google's dialog flow does)

- is it a state?
- when it is a state?

This is a state if you define a complement state:

$$s_{default} = s^c = S \setminus \{s_{cat}^1, s_{cat}^2, s_{dog}^1\}$$

Then your task is to define a policy which will reach a terminal state: when all slots for all intents are set or proven they couldn't be

Often, action for the default state is fallback to conversational

...Still, is S a well-defined state space



A definition

a is a predefined answer from a bot to a user which makes(convinces) the user to fill some slot or reveal an intent

also, there's hidden 'a' part - change bot state (if needed)

```
intent_1(_, _, 'slot_3_val') ->
```

```
db.save(intent_1.slot_3)
```

```
ask_a_user_fill_slot_1_for_intent_1()
```

why we need policy?

(why plain frames are not sufficient)

ambiguity resolution!(context, pronouns, ...) based on previous observations

No ambiguities, then s as defined above is a state

intent_1("panda", "eats", "shoots" | bank in history) -> call_the_police

intent_1("panda", "eats", "shoots" | animal in history) -> take_a_photo

$$\pi : S \rightarrow A$$

...**for now** your model is fully tractable

SM(recall)

makes you able to define policies

"wanna pizza margarita @ work, pay cash"

o(bservervation) = `order_pizza(loc=_, pay=cash)`

H(istory): work: loc = 49,72

S(tate) = `order_pizza(loc=49,72, pay=cash)`

π (olicy): $S \rightarrow A$ (ction space)

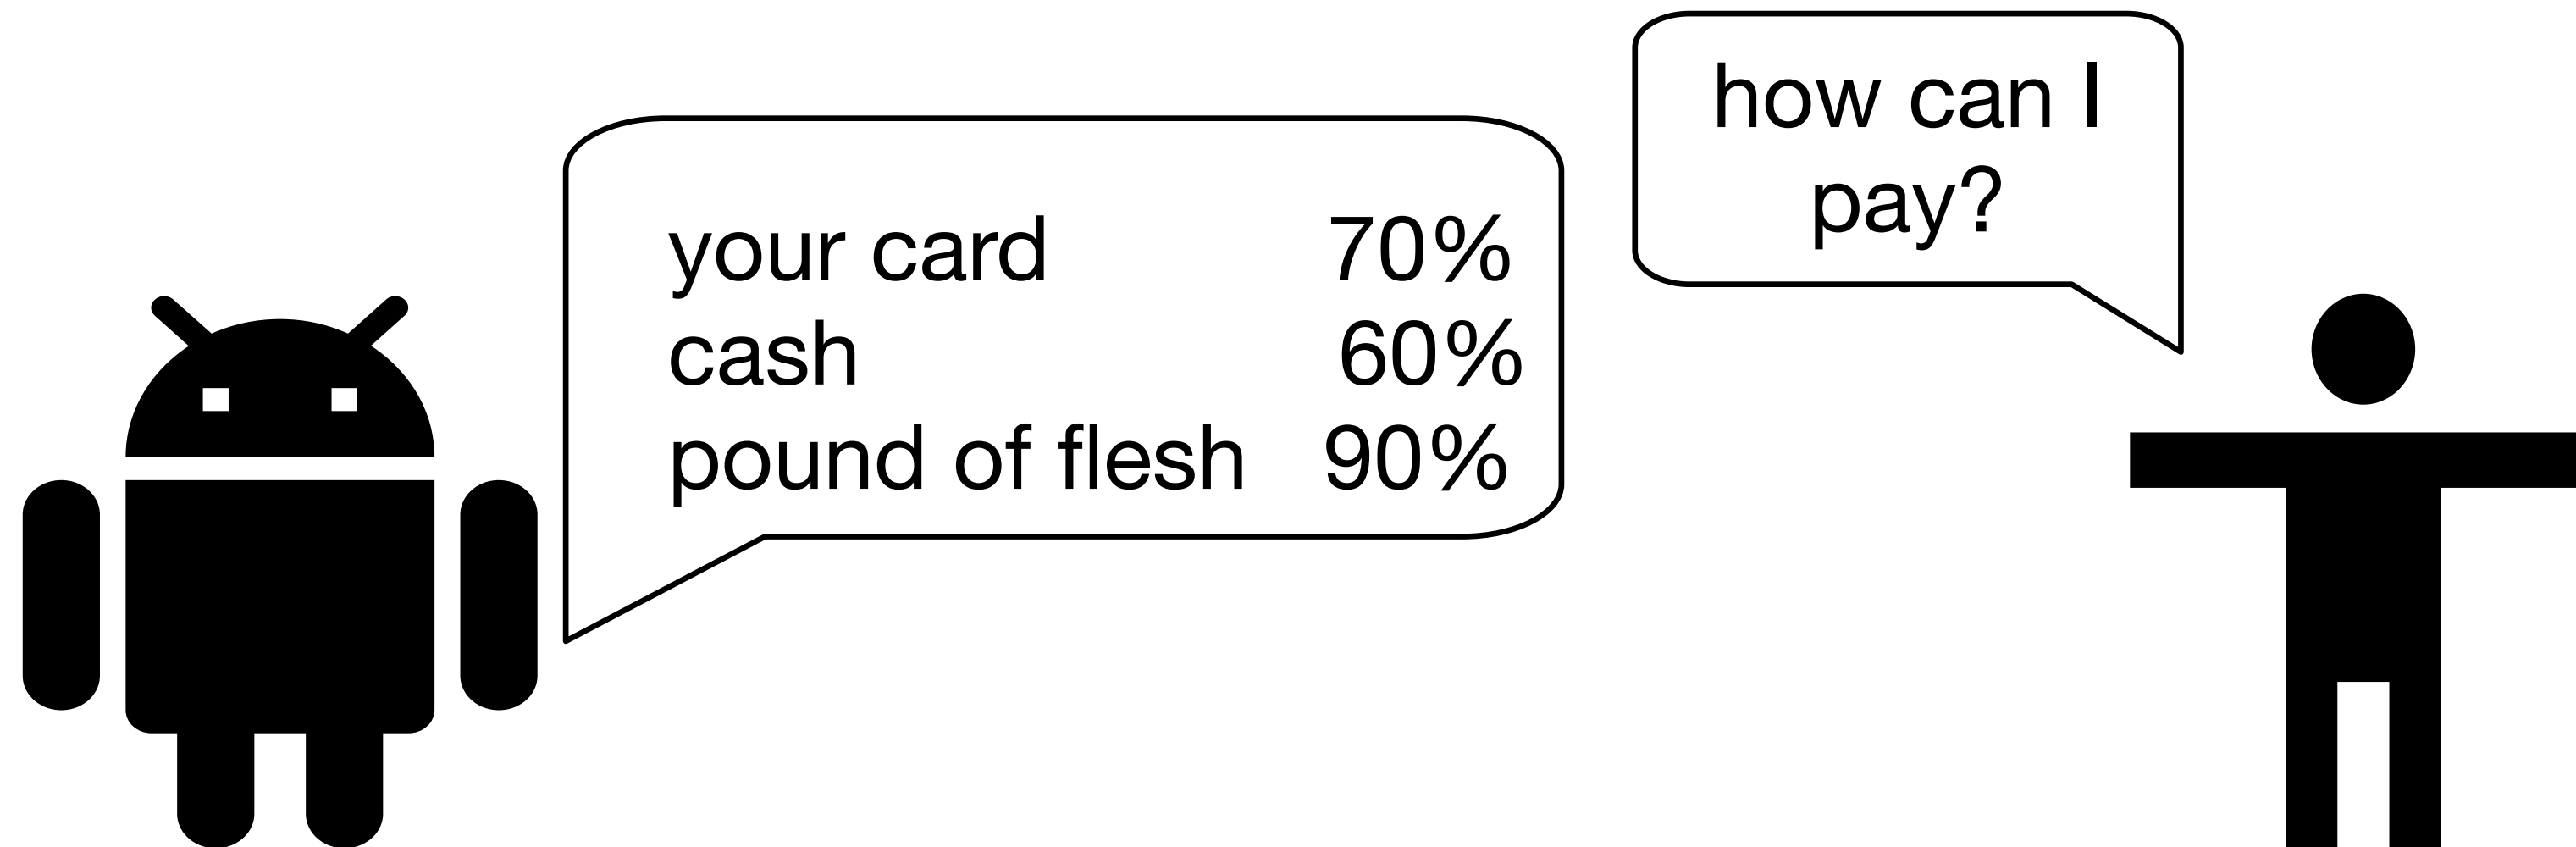
π (`order_pizza(loc=*, pay=*)`) =

= `a_deliver(order(loc=loc, pay=pay))` = `a(ction)`

same works for ranging

text query(TQ)

the best text search you can
by documents you have



- + pretrained models available
- hard/impossible train on own data
- hardcore labeling

BiDAF

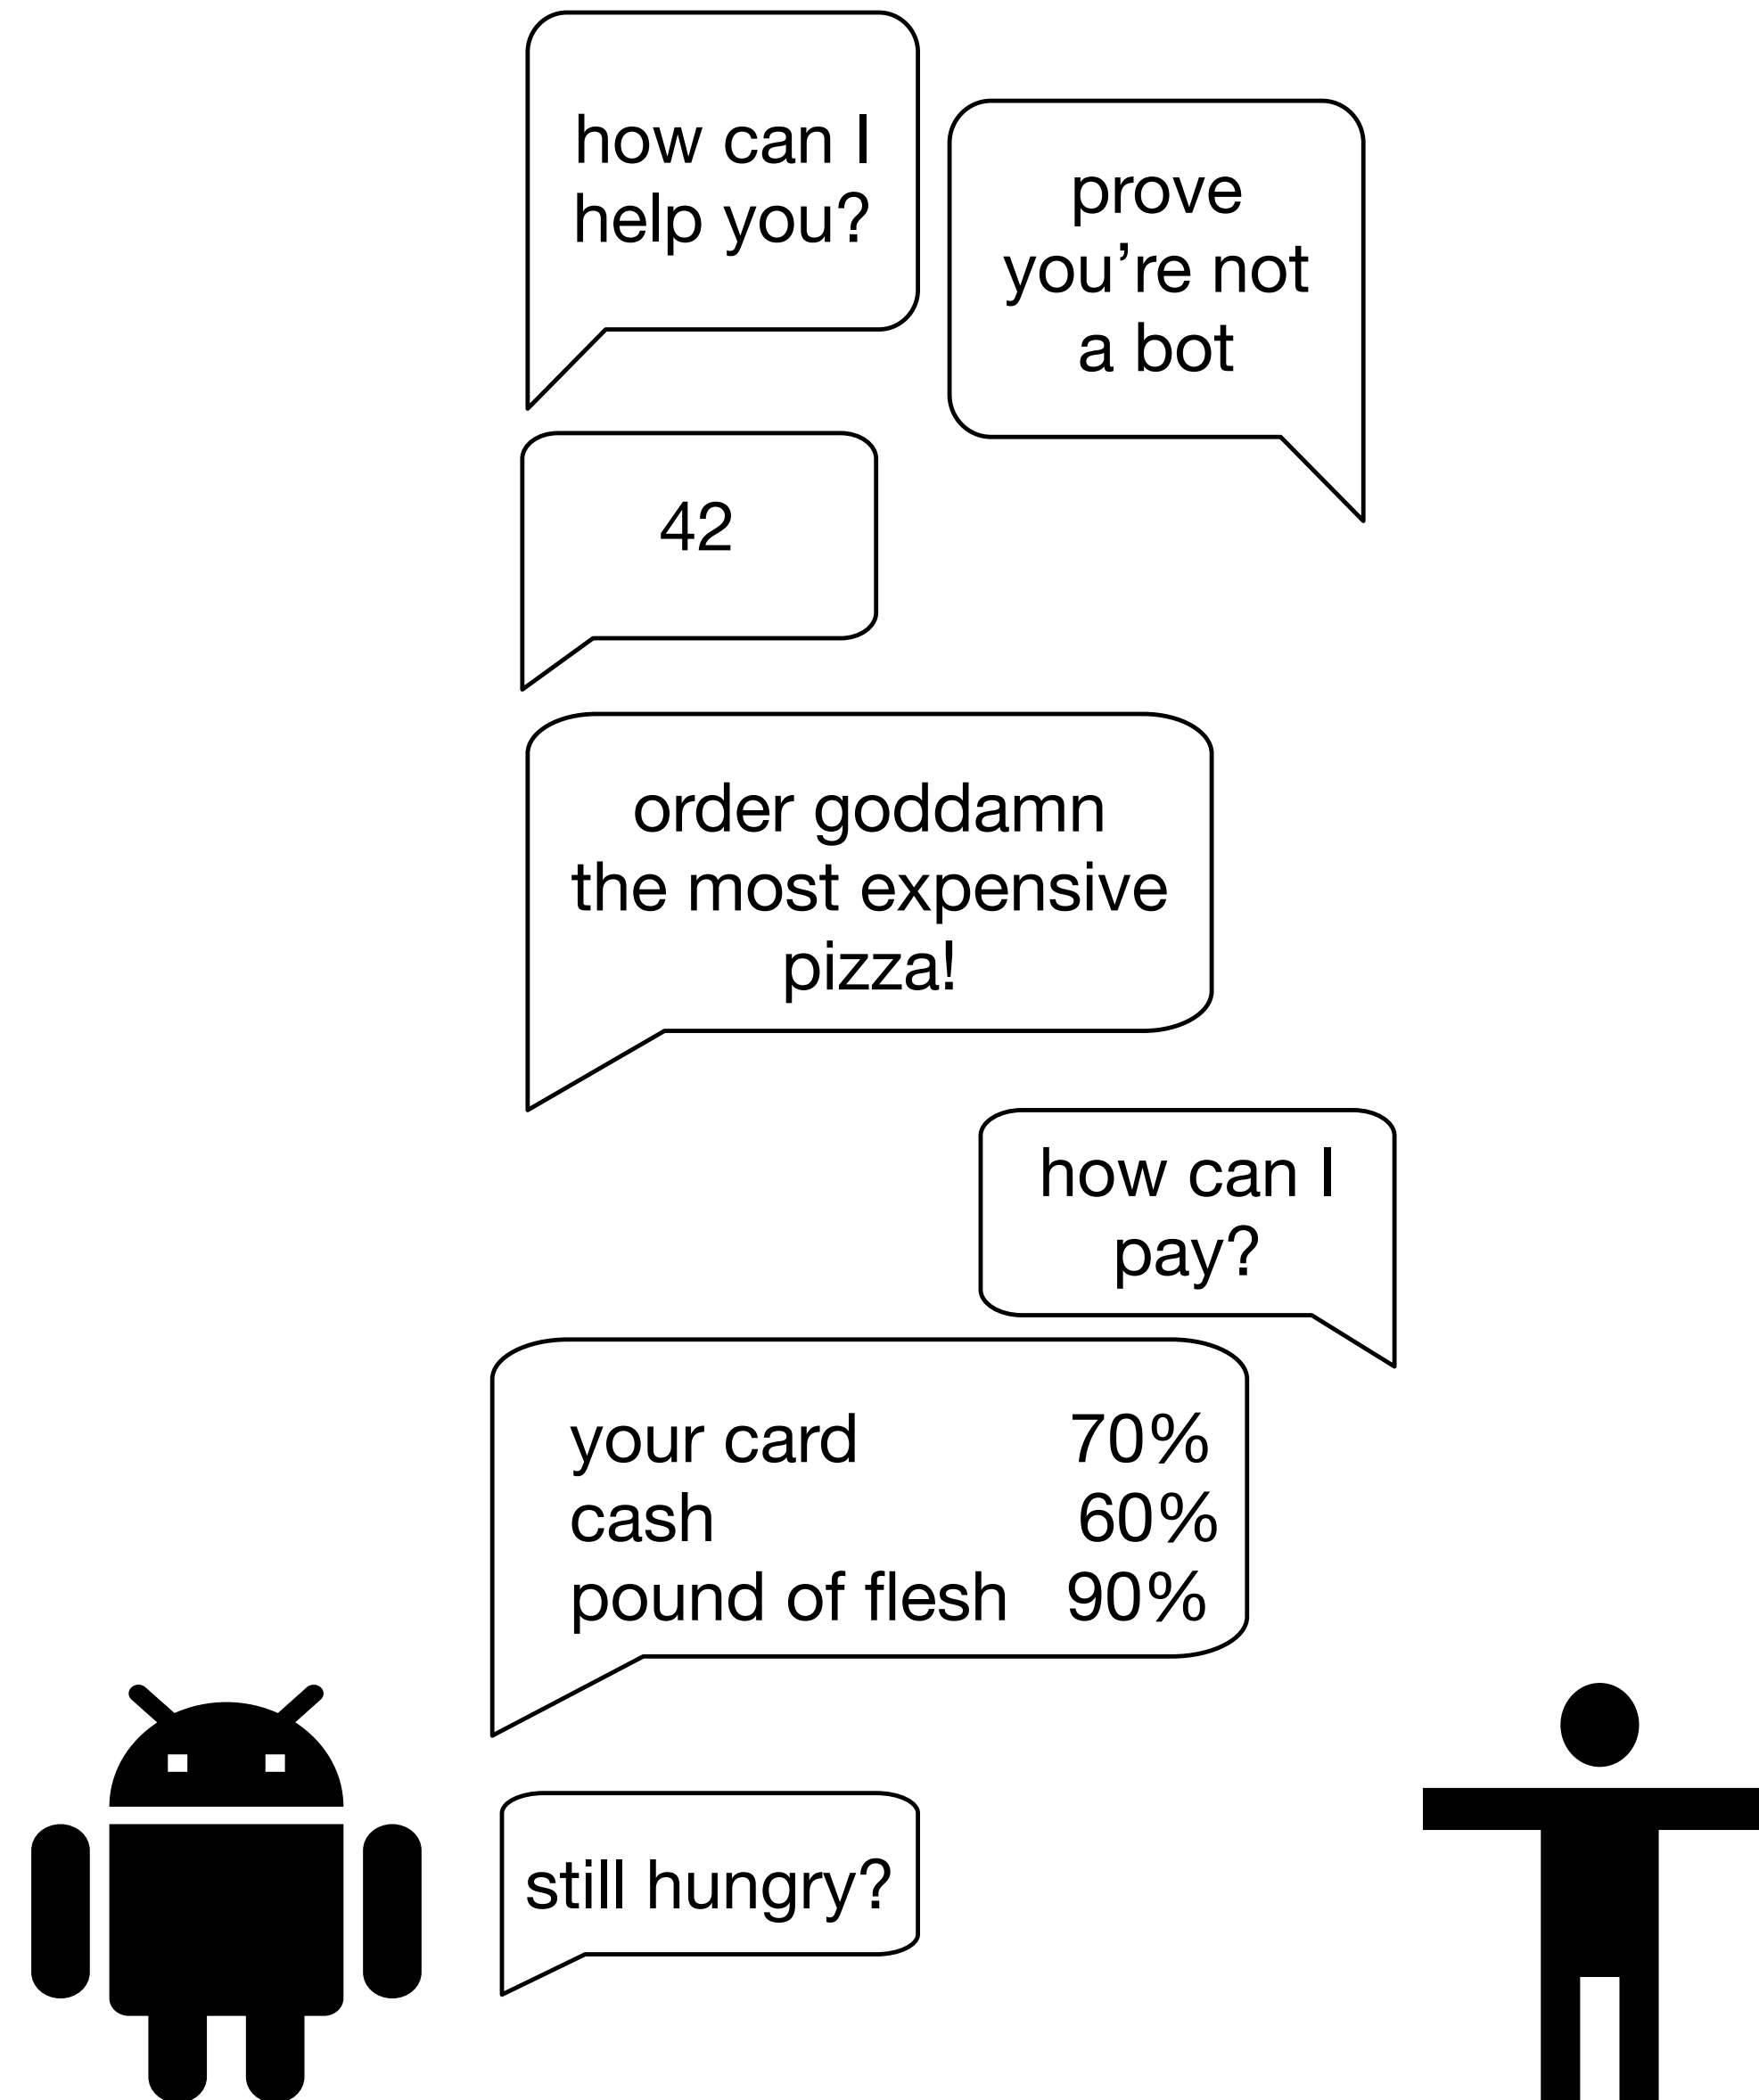
<https://allenai.github.io/bi-att-flow/>

bAbi

<https://research.fb.com/downloads/babi/>

Used if you have some unstructured text DBs, like manuals, HOWTOs, etc.

All 3 within a dialogue



company size



ensembles
ranging
 \mathbb{R} RL

TQ

s2s+att

s2s

forms

att

log scale

data

Ecosystem?

Yandex Алиса

https://www.youtube.com/watch?v=law_tey00Q

Amazon Alexa + skills

Google dialog flow

<https://dialogflow.com/>

Storyline(Alexa skills)

<https://getstoryline.com/>

Deep Pavlov

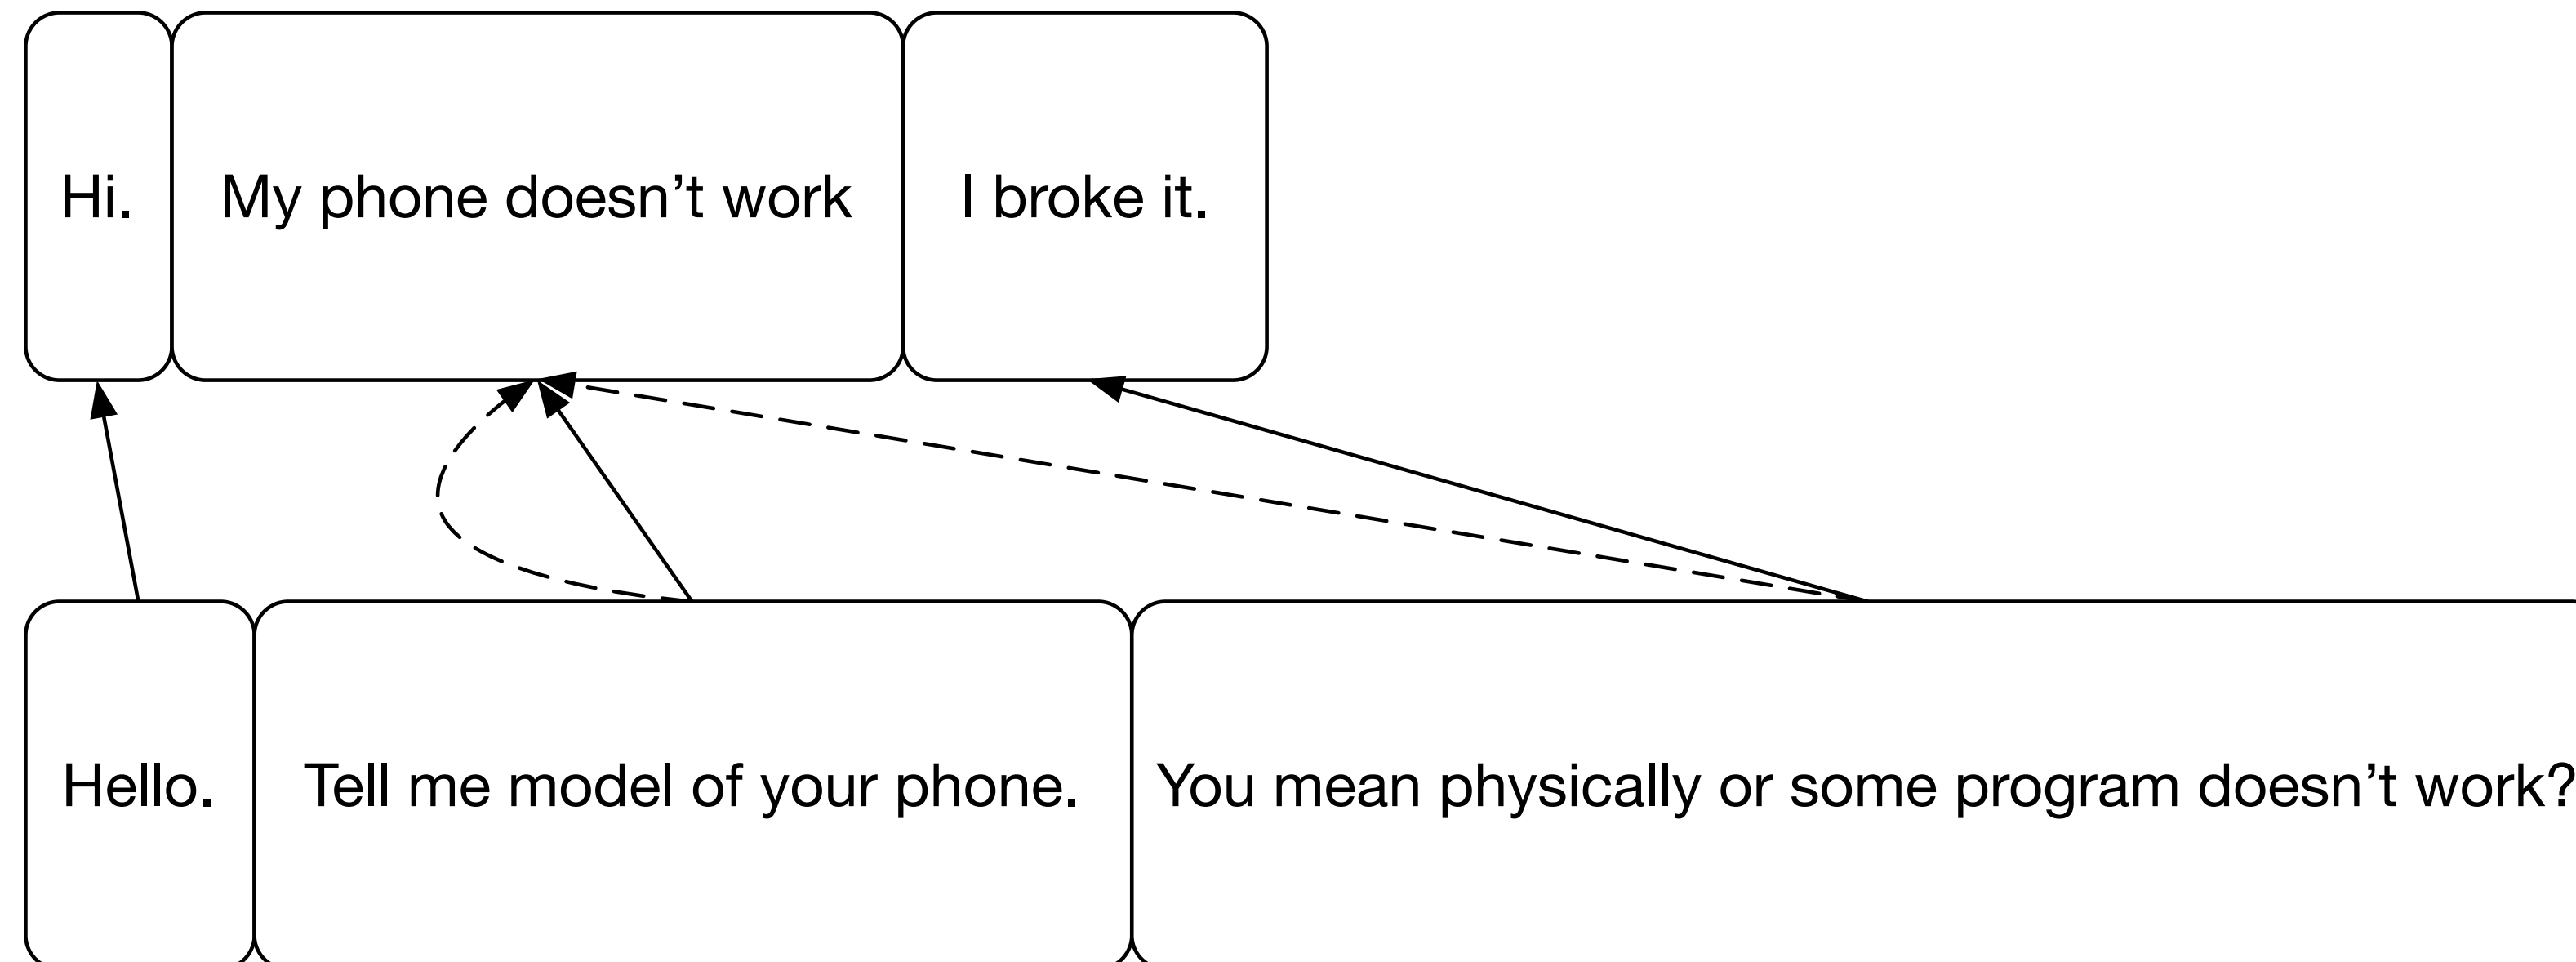
<https://deeppavlov.ai/>

some open problems

Chinese room (open domain)



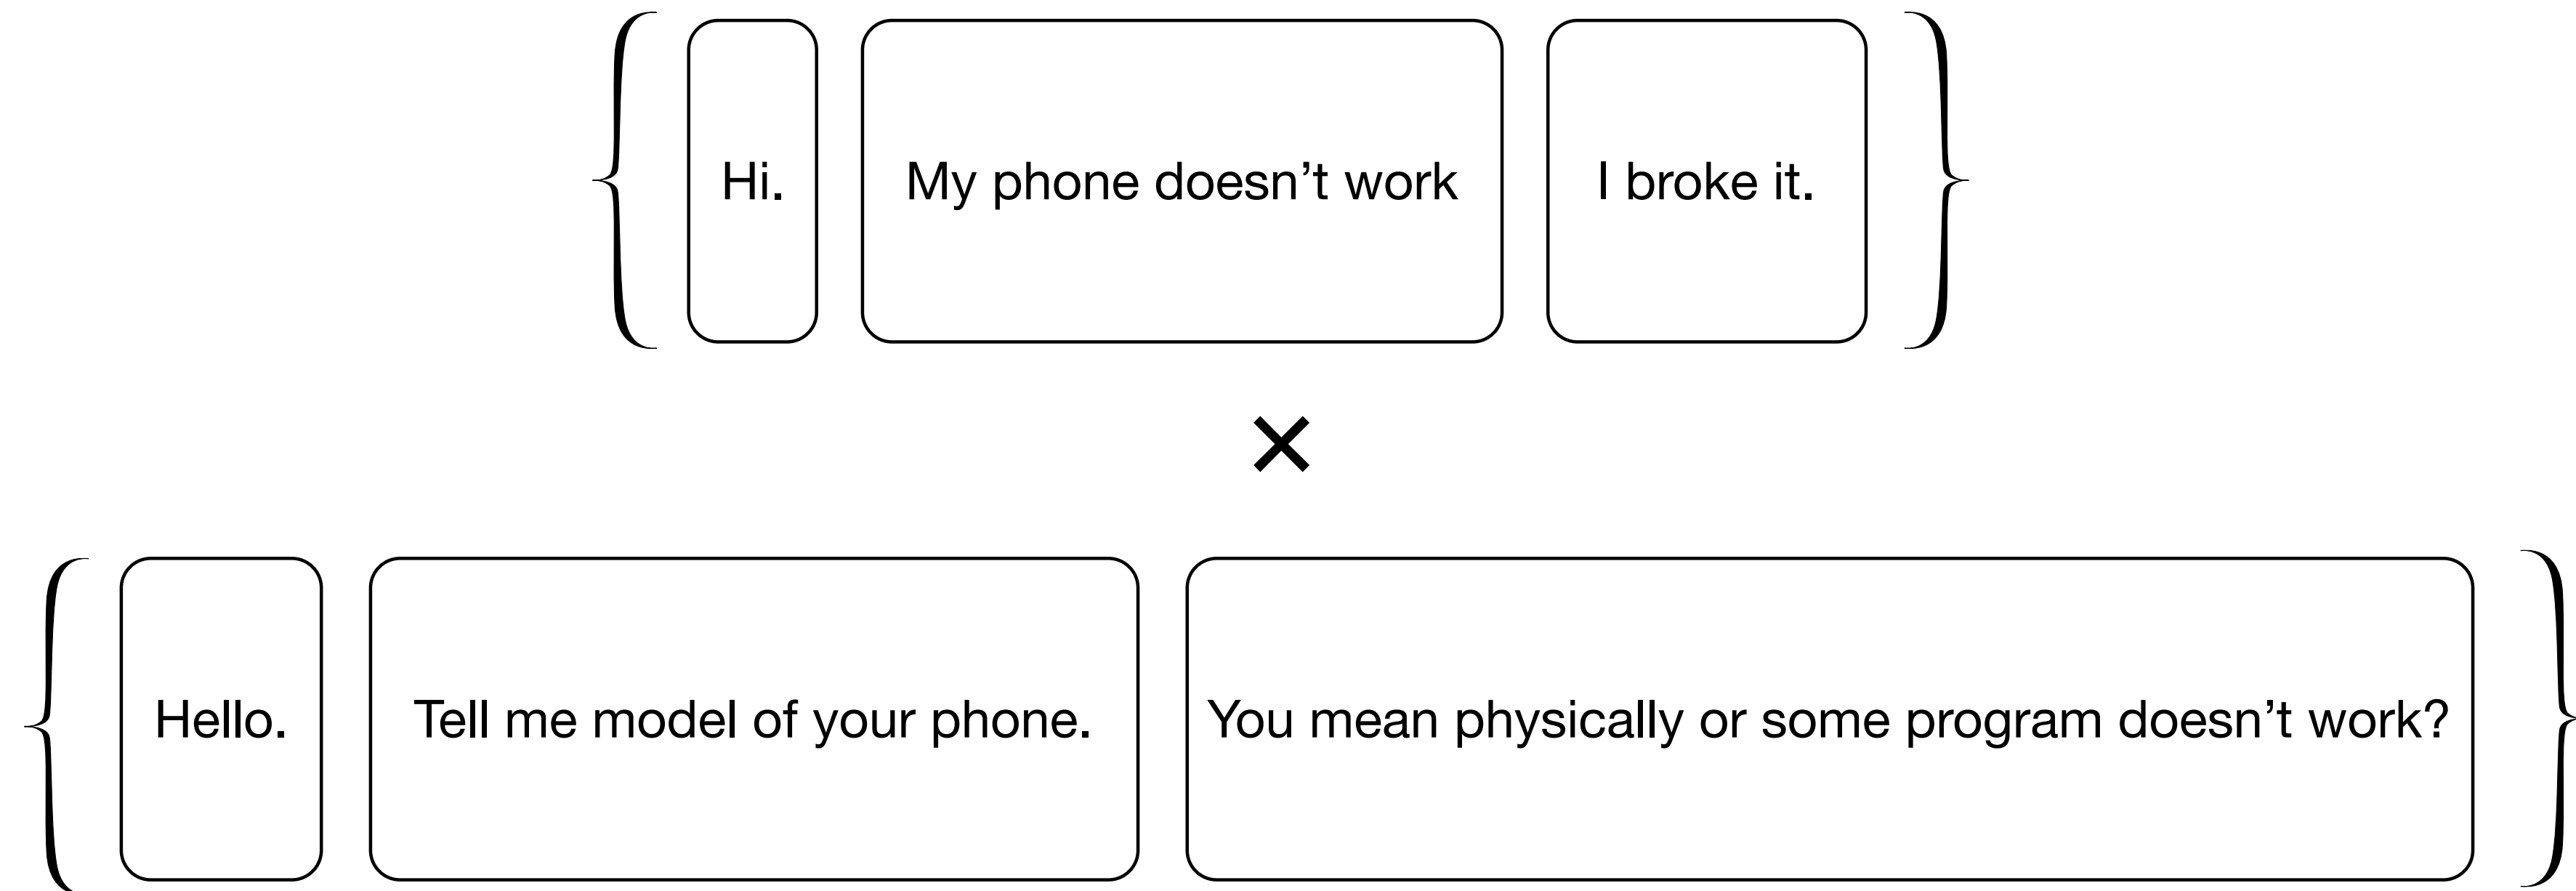
utterance-utterance



Either boring data labeling...

... or merge utterance in a single sentence?

too long | high I/O variation sentences stuck in "I don't know"



then nmt on I/O all pairs

sounds good, does work. But don't know why.

chat context

problem only for conversational

for sentence embedding - you can simply add avg
so far to each embedding as a context

<https://ai.googleblog.com/2018/05/smart-compose-using-neural-networks-to.html>

persona

**placeholders back
substitution**



thanks!

I'm open for consulting & jobs

m3ucat@gmail.com